

Numerical Optimization

Until now, we have discussed evaluating likelihoods.

- We've seen that in many cases maximizing a likelihood can be challenging and require numerical procedures.

Given that we can compute the value of a likelihood, for a $\vec{\theta}$ and \vec{y} we will now discuss numerical optimization techniques.

- Numerical methods typically won't find the exact MLE.
- We will need to set a tolerance level for the quality of our approximation.

Grid Search

Let our parameter vector $k \vec{\theta} \in \mathbb{R}^k$. We define a k -dimensional "grid" or hypercube of points, \mathbb{H}

- We can define a univariate grid, $\mathbb{H}^{(w)}$ for $\theta_1 \in \vec{\theta} \cdot \{\theta_{i,1}, \theta_{i,2}, \dots, \theta_{i,m_i}\}$
- $\mathbb{H} = \mathbb{H}^{(1)} x \mathbb{H}^{(2)} x \dots x \mathbb{H}^{(k)}$
- Ex: $N(\mu, \sigma^2)$

$$\begin{array}{ccccccc} -5 & . & \mu & . & 5 & & \\ . & & . & & . & & \\ . & . & . & . & . & & \\ \sigma^2 & . & . & . & . & & \\ . & . & . & . & . & & \\ . & . & . & . & . & & \end{array}$$

- Often such a grid search is simply equally spaced, but this is not required

Ex. AR(1)

Suppose $c = 0$ and $\sigma^2 = 1$. $\vec{\theta} = \phi \implies k = 1$

- The grid can be a set of values equally spaced between -0.99 and 0.99
- Given a data set \vec{y} , we can compute the exact or conditional likelihood for each ϕ_i in the grid and pick the one that yields the highest value of the log likelihood: ϕ^* .
- Refine the grid search around ϕ^* until our tolerance is reached.

If the likelihood is concave, then we can use a more efficient algorithm:

- Binary search
 1. Pick two adjacent points, θ_j and θ_{j+1} , in the middle of the grid and evaluate the likelihood

2. If $\mathcal{L}(\theta_{j+1}) < \mathcal{L}(\theta_j)$, set the lower bound of the grid to be θ_j and otherwise set the upper bound to be θ_{j+1} .
3. Return to step 1, until the lower and upper bounds of the grid are separated by no more than one point.

$$\begin{array}{ccccccccccc}
 . & . & . & . & . & . & . & . & . & . & . \\
 & & & & \uparrow & & & & \uparrow & & \\
 & & & & L & & & & u & &
 \end{array}$$

- Golden Search: see Heer and Maussner

Grid search is very ineffective when k is large because the number of grid points grows exponentially: doubling the number of points in each dimension results in 2^k as many points.

- This is called the curse of dimensionality.

Newton-Raphson

Define:

$$\vec{g}(\vec{\theta}) = \nabla \ell(\vec{\theta}) = \frac{\gamma \ell(\vec{\theta})}{\gamma \vec{\theta}}$$

$$H(\vec{\theta}) = \nabla^2 \ell(\vec{\theta}) = \nabla \vec{g}(\vec{\theta}) = \frac{\gamma^2 \ell(\vec{\theta})}{\gamma \vec{\theta} \vec{\theta}'}$$

Suppose $H(\vec{\theta})$ is positive definite:

$$\vec{x}' H(\vec{\theta}) \vec{x} > 0 \quad \vec{x} \in \mathbb{R}^k$$

We can approximate $\ell H(\vec{\theta})$ with a 2nd order Taylor expansion around $\vec{\theta}^{(0)}$:

$$\tilde{\ell}(\vec{\theta}) = \ell(\vec{\theta}^{(0)}) + \vec{g}(\vec{\theta}^{(0)})'(\vec{\theta} - \vec{\theta}^{(0)}) + \frac{1}{2}(\vec{\theta} - \vec{\theta}^{(0)})' H(\vec{\theta}^{(0)}) (\vec{\theta} - \vec{\theta}^{(0)})$$

The Newton-Raphson method chooses $\vec{\theta}^{(1)}$ to maximize $\tilde{\ell}(\vec{\theta})$:

$$\left. \frac{\gamma \tilde{\ell}(\vec{\theta})}{\gamma \vec{\theta}} \right|_{\vec{\theta}=\vec{\theta}^{(1)}} = \vec{g}(\vec{\theta}^{(0)}) + H(\vec{\theta}^{(0)}) (\vec{\theta}^{(1)} - \vec{\theta}^{(0)}) = 0$$

$$\implies \vec{\theta}^{(1)} - \vec{\theta}^{(0)} = -H(\vec{\theta}^{(0)})^{-1} \vec{g}(\vec{\theta}^{(0)})$$

$$\implies \vec{\theta}^{(1)} = \vec{\theta}^{(0)} - H(\vec{\theta}^{(0)})^{-1} \vec{g}(\vec{\theta}^{(0)})$$

Newton-Raphson begins with a guess $\vec{\theta}^{(0)}$ and iteratively computes:

$$\vec{\theta}^{(i+1)} = \vec{\theta}^{(i)} - H(\vec{\theta}^{(i)})^{-1} \vec{g}(\vec{\theta}^{(i)})$$

until $\|\vec{\theta}^{(i+1)} - \vec{\theta}^{(i)}\| < \tau$

where τ is some tolerance level.

- Newton-Raphson converges fast if the likelihood is concave and the initial guess is good enough.
- A modified version of NR computes:

$$\vec{\theta}^{(i+1)} = \vec{\theta}^{(i)} - s H(\vec{\theta}^{(i)})^{-1} \vec{g}(\vec{\theta}^{(i)})$$

for various values of s and chooses $\vec{\theta}^{(i+1)}$ that yields the biggest log likelihood.

- Various modified NR methods have been proposed which substitute other positive definite matrices for H^{-1} .

- These are advantageous if H is not possible to compute or invert.
- Typically these are slower but more robust.

If analytical derivatives are not possible, numerical derivatives are an option.

- e.g. the i th element of $\vec{g}(\vec{\theta})$ can be approximated with

$$g_i(\vec{\theta}) = \frac{1}{\Delta} (\ell(\theta_1, \dots, \theta_i + \Delta, \dots, \theta_k) - \ell(\theta_1, \dots, \theta_i, \dots, \theta_k)) \text{ for some small } \Delta.$$

- The Hessian can be computed numerically from \vec{g} in a similar fashion.